Extensions for DDPG and analysis of its components

subtitle here

**Maximilian Gehrke · Tabea Wilke ·
Yannik Frisch**

Abstract TODO

1 Introduction

Deep Deterministic Policy Gradients (DDPG) arises from Deterministic Policy Gradients (DPG) and Deep Q-Learning (DQN). In the following we describe the underlying algorithms DPG and DQN and which aspects DDPG uses of both of them.

1.1 Deep Q-Learning (DQN)

The DQN approach combines the approximation power of Neural Networks with traditional Q-learning. It enables solving the classic Reinforcement Learning problem of achieving the maximum expected reward over time, even for large state spaces (e.g. image frames). The algorithm is an off-policy, model-free approach and is able to find a close to optimal action-value function for many cases [x] and from this a close to optimal deterministic policy by greedily selecting the action: $\pi(s) = \max_a Q^*(s, a)$. In terms of a formula the optimal

F. Author
first address
Tel.: +123-45-678910
Fax: +123-45-678910
E-mail: fauthor@example.com

S. Author
second address

action-value function is represented by

$$Q^*(s_t, a_t) = \max_\pi E\left[\sum_{t'=t}^{T} \gamma^{t'-t} r_{t'} | s_t = s, a_t = a, \pi\right]$$

where $\gamma \in [0, 1)$ is called the discount factor, controlling the agents preference for rewards closer or further away in time. Rewards later on in an episode will still have impact on the result but their influence decreases by the amount of time-steps required to reach them in the future. By definition this optimal value function yields the Bellman equation [x] and can be reinterpreted as maximizing the current reward and the discounted action-value of the resulting state. In formula this gives:

$$Q^*(s, a) = E_{s' \sim \epsilon}\left[r + \gamma \max_{a'} Q^*(s', a') | s, a\right]$$

For approximating the action-value function $Q(s, a|\theta) \approx Q^*(s, a)$ the approach uses a deep neural network, called the Q-Network. The Q-Network can be trained by minimizing a sequence of loss functions $L_i(\theta_i)$, depending on it's weights:

$$L_i(\theta_i) = E_{s, a \sim \rho(.), s' \sim \epsilon}\left[\left(r + \gamma \max_{a'} Q(s', a'|\theta_{i-1}) - Q(s, a|\theta_i)\right)^2\right]$$

This loss function is similar to the classical temporal-difference loss used in Q-Learning, but with approximated action-value functions instead of lookup-tables. Derivating this loss w.r.t. the approximation's weights gives:

$$\nabla_{\theta_i} L_i(\theta_i) = E_{s, a \sim \rho(.), s' \sim \epsilon}\left[\left(r + \gamma \max_{a'} Q(s', a'|\theta_{i-1}) - Q(s, a|\theta_i)\right) \nabla_{\theta_i} Q(s, a|\theta_i)\right]$$

This gradient can be used to optimize the loss function by using stochastic gradient descent.

Furthermore, a replay buffer is used which stores samples of the environment. This allows random mini-batch sampling, which decorrelates the samples and is proven to improve the performance by greater data efficiency [x]. The mini-batch sampling also enables the use of improved derivatives of vanilla stochastic gradient descent, e.g. RPROP as in the 'neural fitted Q-Learning (NFQ)' [x] approach or ADAM update [x].

There are different ways of estimating the expected Q-values. Either with a target network with the same structure as the network for the action-value function or the normal network. If a target network is used, the target weights need to be updated after some training steps.

A pseudocode for the DQN approach can be found in 1. The DQN approach was able to significantly outperform earlier learning methods despite incorporating almost no prior knowledge about the inputs [x], but is limited by the disabilty to cope with continous and high-dimensional action spaces due to the max operator in the action selection [x]. This limitations can be adressed by combining the approach with the Deterministic Policy Gradient, which is described in the following section.

Algorithm 1 Deep Q-Learning (DQN)

Initialize: Replay buffer D with high capacity
Initialize: Neural network for action-value function Q with random weights θ
Initialize: Neural network for target action-value function $\hat{Q}$ with weights $\theta^- = \theta$
  for episode 1 **to** M **do**
    reset environment to state s_1
    for $t = 1$ **to** T **do**
      if random $i \leq \epsilon$ **then**
        random action a_t
      else
        $a_t = \mathrm{argmin}_a Q(s_t, a | \theta)$
      end if
      execute $a_t \rightarrow$ reward r_t and next state s_{t+1}
      save (s_t, a_t, r_t, s_{t+1}) in D
      sample minibatch (s_i, a_i, r_i, s_{i+1}) from D
$$q_i = \begin{cases} r_i & \text{if episode terminates at step } i+1 \\ r_i + \gamma \max_{a'} \hat{Q}(s_{i+1}, a' | \theta^-) & \text{else} \end{cases}$$
      perform gradient descent on $(q_i - Q(s_i, a_i; \theta))^2_\theta$
      every C steps update $\hat{Q} = Q$
    end for
  end for

1.2 Deterministic Policy Gradient (DPG)

The most problems in reinforcement learning have an continuous action space, which makes it very difficult to choose the best action given a policy. In a stochastic point of view, there is a probability distribution which represents the policy. The policy is obtained through $\pi_\theta(a|s) = P[a|s; \theta]$. The goal is to achieve the best possible return and therefore choose the best action by solving the problem with a gradient over the total reward. This is only possible through solving an integral over all actions and actions. In the deterministic view the policy is a discrete mapping from state to actions and thus only one integration over the state space is sufficient. The integration is done via importance sampling. As a result of this, stochastic policy gradients need much more samples than deterministic policy gradients. To handle the exploration-exploitation dilemma, the idea of Silver et al. is to use a stochastic behaviour policy and a deterministic target policy. The behaviour policy should ensure that the exploration is big enough and the target policy should exploit enough based on the policy gradient. The algorithm they used was an off-policy actor-critic method. Thereby, the action-value function is estimated with a function approximation. The policy parameters will be updated in direction of the gradient of the action-value function. The direction of the gradient underlies the *policy gradient theorem*:

$$\nabla_\theta J(\pi_\theta) = \int_S \rho^\pi(s) \int_A \nabla_\theta \pi_\theta(a|s) Q^\pi(s, a) \mathrm{d}a \mathrm{d}s$$
$$= E_{s \sim \rho^\pi,\, a \sim \pi_\theta}[\nabla_\theta \log \pi_\theta(a|s) Q^\pi(s, a)]$$

1.3 Actor-Critic Methods

The advantage of actor-critic methods is that they learn policies as well as value functions. To get an intuition about these methods the following figure illustrates the update-cycle:

Fig. 1 Intuition about actor-critic methods (figure from [5])

The actor is responsible for the change of the policies, the critic has to update the parameter of the state-value function. Updating the actors' and the critics' parameter follows the TD-error of the critic which is produced through the reward and the current error of the estimated state values. As Fig. 1 illustrates, the actor has no information about the current reward and the critic has no direct influence on the actions. The update of actor and critic can be formulated as follows:

Fig. 2 Intuition about actor-critic methods (figure from [5])

1.4 Deep Deterministic Policy Gradients (DDPG) [1]

DDPG is a model-free and off-policy algorithm which arises from the combination of DQN and DPG. As a result of a continuous action space which leads to a not possible application of Q-learning. Therefore, it is a policy gradient algorithm which uses actor-critic methods with a deterministic target policy and deep Q-learning. Both the actor and the critic will be done with neural networks to represent the parameters. The actor will choose an action from a continuous action space.

- DQN uses deep networks to estimate the action-value function
  - it can only handle discrete and low-dim action spaces
- discretizing the action space often suffers from the course of dimensionality
- PolicyGradientTheorem from continous space to discrete space presented in DPG paper
- naive extension of DPG with nns turns out to be unstable for challenging problems
- Deep DPG (DDPG): combination of DQN and DPG, where:
  - networks are trained off-policy with samples from a replay buffer to minimize the temporal correlations between samples
  - the networks are trained with target networks to give consistent targets during temporal difference backups
  - batch normalization is used
- DDPG is able to learn from low dim observations (torques etc.), aswell as from high dim observations in pixel space

2 Extensions to the Algorithm

References

1. Lillicrap TP, Hunt JJ, Pritzel A, Heess N, Erez T, Tassa Y, Silver D, Wierstra D (2015) Continuous control with deep reinforcement learning. arXiv preprint arXiv:150902971
2. Mnih V, Kavukcuoglu K, Silver D, Graves A, Antonoglou I, Wierstra D, Riedmiller M (2013) Playing atari with deep reinforcement learning. arXiv preprint arXiv:13125602
3. Mnih V, Kavukcuoglu K, Silver D, Rusu AA, Veness J, Bellemare MG, Graves A, Riedmiller M, Fidjeland AK, Ostrovski G, et al. (2015) Human-level control through deep reinforcement learning. Nature 518(7540):529
4. Silver D, Lever G, Heess N, Degris T, Wierstra D, Riedmiller M (2014) Deterministic policy gradient algorithms. In: ICML
5. Sutton RS, Barto AG (2018) Reinforcement learning: An introduction. MIT press