

Reviewing Natural Actor Critic methods

Maximilian A. Gehrke ·
Yannik P. Frisch · Tabea A. Wilke

Received: date / Accepted: date

Abstract In this paper we describe the natural actor critic approach and provide an extensive overview about the current research. This includes a basic description of the natural gradient, actor critic approaches and comparisons between existing extensions. Additionally, we improve the episodic Natural Actor Critic algorithm by applying it with two neural networks instead of basis functions.

Keywords Natural Gradient · Neural Networks · Advantage Function · Actor Critic · NAC

1 Introduction

Natural actor critic (NAC) methods [26] have been very successful in the last years. They could be applied to various fields, including traffic optimization [27], dialog systems [14] and high dimensional control tasks [22, 23, 24, 25]. NAC methods belong to the policy gradient family with the characteristics of employing the natural gradient, an actor critic approach and a compatible function approximation. Natural gradient algorithms have been applied successfully in various fields such as road traffic optimization [27], robotic control tasks [16], motor primitive learning [22] and locomotion of a two-linked robot arm [20]. Using the natural gradient has several advantages and properties, which are listed in the next section.

Maximilian A. Gehrke
E-mail: maximilian_alexander.gehrke@stud.tu-darmstadt.de

Yannik P. Frisch
E-mail: yannik_phil.frisch@stud.tu-darmstadt.de

Tabea A. Wilke
E-mail: tabeaalina.wilke@stud.tu-darmstadt.de

Every gradient algorithm has very good sample complexity guarantees [19]. Policy gradient methods represent the policy using differentiable function approximation. They optimize a scalar performance measure $J(\theta)$, called cost function by repeatedly estimating it's gradient w.r.t. the policy parameters θ and updating the policy parameters θ a proportion in that direction: $\theta_{t+1} = \theta_t + \alpha \nabla J(\theta_t)$. PGM have several advantages. They are model-free, have good convergence properties, can learn stochastic policies and are effective in high-dimensional or continuous action spaces. However, PGM is typically inefficient, has high variance and it typically converge to a local rather than a global optimum. Further, with PGM we can introduce a prior on the policy, we can converge to a deterministic policy and don't have to chose a suboptimal action for exploration purposes and we can chose our actions stochastically. However, the policy needs to be differentiable w.r.t it's parameters.

Actor-critic methods approximate a value function beside the policy. This means that we have to estimate two set of weights, w for the value function $v_w(s)$ (which is a state-action value function in most cases) and θ for the policy $\pi_\theta(a|s)$. The value function estimator is used to learn the policy, where actor refers to the policy and critic to the value function. The actor tells the agent which actions to execute, the critic rates the observations, updates it's own parameters and finally the actor updates it's parameters w.r.t the critic.

The paper structures itself in the following way: We start by setting up some preliminaries in section 2. In section 3 we introduce the natural gradient and discuss it's properties in section ???. The natural actor critic algorithm will be presented in section 4 and the modifications and extensions currently known in section 5. Finally, we close with a discussion in section 6.

2 Preliminaries

We consider a standard reinforcement learning framework, in which a learning agent interacts with a Markov Decision Process (MDP) [12, 29]. For each discrete time step $t \in \{0, 1, 2, \dots\}$, the state, action and reward is denoted as $s_t \in S$, $a_t \in A$ and $r_{t+1} \in R \subset \mathbb{R}$ respectively. The dynamics of the environment are described by the state-transition probabilities $P_{ss'}^a = \Pr\{S_t = s' | S_{t-1} = s, A_{t-1} = a\}$ and the expected immediate rewards $R_s^a = \mathbb{E}[R_t | S_{t-1} = s, A_{t-1} = a]$, for all $s, s' \in S, a \in A$. The agent's behavior at each time step t is specified by a policy $\pi_\theta(a|s) = \Pr\{A_t = a | S_t = s, \theta\}$, where θ denotes the parameters of the policy.

We assume that π is differentiable w.r.t. it's parameters, so that $\frac{\partial \pi(a|s)}{\partial \theta}$ exists and we can estimate the gradient of the objective function $J(\theta)$ by applying the policy gradient theorem [30]

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(a|s) Q^{\pi_\theta}(s, a)], \quad (1)$$

where $Q^{\pi_\theta}(s, a)$ denotes an action-value function. One of the most basic policy gradient algorithms, *REINFORCE* [34], estimates the action-value function Q^{π_θ} by using the expected discounted return, also known as Monte-Carlo return

$$Q^{\pi_\theta}(s, a) \approx G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}, \quad (2)$$

where γ is a discount factor with $\gamma \in [0, 1]$. Policy gradient methods use the gradient of the objective function $J(\theta)$ and a learning rate $\alpha \in [0, 1]$ to recursively update the parameters of $\pi_\theta(a|s)$, $\theta_{t+1} = \theta_t + \alpha \nabla_\theta J(\theta)$, and find a local optimum.

However, instead of using a Monte-Carlo estimate directly, Actor-Critic methods model the action-value function with a function approximator $Q^{\pi_\theta}(s, a) \approx Q_w(s, a)$ [30]. $Q_w(s, a)$ is called the critic and introduces a second set of parameters, w , which need to be optimized; $\pi_\theta(a|s)$ is called the actor. By introducing a baseline $B(s, a)$, we can reduce the variance of the action-value function estimate and accelerate learning [29]:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(a|s) Q_w(s, a) - B(s, a)]. \quad (3)$$

A good baseline with minimal variance is the value function. Subtracting the value function from the action-value function yields the Advantage function $A(s, a) = Q(s, a) - V(s)$. However, the critic can directly estimate the advantage function $A_w(s, a)$ for computing the gradient:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(a|s) A_w(s, a)]. \quad (4)$$

Optimizing the objective function with vanilla gradient descent is sensitive to parametrization and can be inefficient. One could instead use the *natural gradient*, which is described in the next section.

3 Natural Gradient

[1] The natural gradient was first introduced by Amari in 1998 [3]. The difference between the natural gradient and the ordinary vanilla gradient, is the direction it points to. The ordinary vanilla gradient only points to the steepest direction, if the parameter space has an Euclidean character [x]. The natural gradient, however, points to the steepest direction of a Riemann parameter space (e.g. neural networks [3]).

A Riemannian parameter space is a differentiable manifold, where for each tangent space an inner product $\langle \cdot, \cdot \rangle$ exists. For two tangent vectors \mathbf{u} and \mathbf{v} , the inner product $\langle \mathbf{u}, \mathbf{v} \rangle$ yields a real number. This makes it possible to define notions such as length, areas, angles or volumes. To calculate the gradient, we need to be able to calculate the squared length of a small incremental vector $d\mathbf{w}$ connecting a point \mathbf{w} and $\mathbf{w} + d\mathbf{w}$. Equation 5 shows on the left the

formular for Riemanian spaces and on the right the formular for Euqulidean spaces:

$$|d\mathbf{w}|^2 = \sum_{i,j} g_{ij}(\mathbf{w}) dw_i dw_j, \quad |d\mathbf{w}|^2 = \sum_{i=1}^n (dw_i)^2, \quad (5)$$

where $g_{ij}(\mathbf{w})$ is a function, enabling us to create a measure of distance. It is also written as an $n \times n$ matrix $G = (g_{ij})$, called Riemannian metric tensor, and reduces to the unit matrix I in the case of an Euclidean orthonormal parameter space. Therefore, the riemannian case is a generalization of the Euclidean orthonormal case [10, 3]. We can utilize the Riemannian metric tensor to construct a gradient which points in the steepest direction of Riemannian spaces:

$$\tilde{\nabla}_{\theta} J(\theta) = G^{-1} \nabla_{\theta} J(\theta). \quad (6)$$

$\tilde{\nabla}_{\theta}$ is the natural gradient w.r.t the parameters θ . Learning should be carried out with a gradient descent like update rule: $\theta_{t+1} = \theta_t + \alpha \tilde{\nabla}_{\theta} J(\theta)$. In the special case that the parameter space is Euclidean and the coordinate system is orthonormal, the conventional gradient equals the natural gradient: $\tilde{\nabla}_{\theta} J(\theta) = \nabla_{\theta}$. If the Fisher information matrix (FIM) exists, it could be shown that we can use it in equation 6 as the Riemannian metric tensor and get a natural gradient [23, 2]. The FIM of a policy π_{θ} is defined as:

$$F_{\theta} = \mathbb{E}_{s,a} [\nabla_{\theta} \log \pi_{\theta}(a|s)^T \nabla_{\theta} \log \pi_{\theta}(a|s)]. \quad (7)$$

If we look at the problem from a different angle, we can see the uniqueness of the natural gradient: it's invariance to parameterization [21, 23]. Equation 8 defines our objective. With policy gradient methods, we want to change the parameters of the policy, so that the resulting parameters maximize the objective function J . This can be done by taking the vanilla gradient. However, the vanilla gradient has the downfall that in flat regions of the parameter space, the algorithm will move very slowly, whereas in steep regions the algorithm will move very fast and even shoot beyond the local maximum. This is because for every parametrization θ the gradient is different. This is why we ought to search for a way to measure the distance between two distributions. One of these is the Kullback-Leibler divergence which can be approximated by the second-order Taylor expansion, as shown in equation 9. We constrain the Kullback-Leibler divergence to be less than a fixed value ϵ . This means, that the parameters change exactly for a given distance in the parameterspace. Equation 8 and 9 together form an optimization problem:

$$\max_{\delta\theta} J(\theta + \delta\theta) \approx J(\theta) + \delta\theta^T \nabla_{\theta} J(\theta) \quad (8)$$

$$\text{s.t. } \epsilon = D_{KL}(\pi_{\theta} || \pi_{\theta+\delta\theta}) \approx \frac{1}{2} \delta\theta^T F_{\theta} \delta\theta \quad (9)$$

which solution yields equation 6, applied with the FIM F_{θ} .

- **Online Learning:** The NG can be used online and therefore can learn from incomplete sequences and reduce the variance of the action-value function estimation [21, 23].
- **1st order method:** the natural gradient is a first order method, but implements second order advantages [21]. This is especially relevant for problems, where the cost function is accessible indirectly [9].
- **Better & faster convergence:** In many cases the NG converge faster than vanilla gradient algorithms [28, 3] and avoids getting stuck in plateaus [2, 3].
- **Drawbacks:** The Riemannian metric tensor needs to be nonsingular and invertible. This is not always the case and even if, the inversion of a matrix is very costly. In addition, by applying the NG, the policy variance might reduce to zero. This poses a problem, which is dealt with in the TRPO & PPO algorithms.

4 Natural Actor Critic

In this section we describe the *Natural Actor Critic* (NAC) algorithm [26]. We focus on the trajectory based formulation, called episodic NAC, and present the pseudo code in algorithm 1.

In episodic NAC, we have a fixed amount of updates u and a fixed amount of steps the agent executes in the environment every update. Therefore, if a trajectory e has reached a terminal state before the agent executed all its steps, the algorithm samples a new trajectory. This repeats until the maximum number of steps is met and the current trajectory is interrupted. During this process, all states we see, all actions we take and all rewards we get are stored for each trajectory.

After sampling, we perform the critic evaluation. We determine the compatible function approximation, the basis functions and the reward statistics for the samples of a single episode and solve a linear equation system to get w_e . We update w_t by adding w_e multiplied by a learning rate $\beta \in [0, 1]$. We repeat this process for all trajectories encountered during the update. Then, we check if the angle between w_{t+1} and w is smaller than some fixed value ϵ and if so, we update the policy parameters θ by adding w_{t+1} multiplied by a learning rate $\alpha \in [0, 1]$.

Our cost function is the expected return $\mathbb{E}[G_t]$, which can be written as the discounted sum of advantages, which in return can be written in terms of the expected reward and value function [26]

$$\sum_{t=0}^N \gamma^t A(s_t, a_t) = \sum_{t=0}^N \gamma^t r(s_t, a_t) + \gamma^N V(S_{N+1}) - V(S_0), \quad (10)$$

Algorithm 1 Episodic Natural Actor Critic (eNAC)

Require: Parameterized policy $\pi_\theta(a|s)$ and it's derivative $\nabla_\theta \log \pi_\theta(a|s)$ with initial parameters $\theta = \theta_0$.

- 1: **for** $u = 1, 2, 3, \dots$ **do**
- 2: **for** $e = 1, 2, 3, \dots$ **do**
- 3: **Execute roll-out:** Draw initial state $s_0 \sim p(s_0)$
- 4: **for** $t = 1, 2, 3, \dots, N$ **do**
- 5: Draw action $a_t \sim \pi_{\theta_t}(a_t|s_t)$, observe next state $s_{t+1} \sim p(s_{t+1}|s_t, a_t)$ and reward $r_{t+1} = r(s_t, a_t)$.
- 6: **end for**
- 7: **end for**
- 8: **Critic Evaluation (repeat for each sampled trajectory):** Determine compatible function approximation of advantage function $A(s, a) \approx A_{w_t}(s, a)$.
- 9: Determine basis functions: $\Phi_e = \left[\sum_{t=0}^T \gamma^t \nabla_\theta \log \pi_\theta(a_t|s_t)^T, 1 \right]^T$,
reward statistics: $R_e = \sum_{t=0}^T \gamma^t r_t$ and solve $\begin{bmatrix} w_e \\ J \end{bmatrix} = (\Phi_e^T \Phi_e)^{-1} \Phi_e^T R_e$.
Update critic parameters: $w_{t+1} = w_t + \beta w_e$.
- 10: **Actor Update:** When the natural gradient is converged, $\angle(w_{t+1}, w_t) \leq \epsilon$, update the policy parameters: $\theta_{t+1} = \theta_t + \alpha w_{t+1}$.
- 11: **end for**

where N is the number of steps executed in a trajectory. If we assume $\gamma \neq 1$, we can remove the term $\gamma^N V(S_{N+1})$, because in the limit the term becomes zero ($\lim_{N \rightarrow \infty} \gamma^N = 0$). Additionally, if we assume that we always start in the same start state S_0 , we can write $V(S_0)$ as our cost function $J(\theta)$:

$$\sum_{t=0}^N \gamma^t A(s_t, a_t) = \sum_{t=0}^N \gamma^t r(s_t, a_t) - J(\theta). \quad (11)$$

One of the key aspects of the NAC algorithm is the use of a compatible function approximation $A_w(s, a)$ to estimate $A(s, a)$ [30], which by definition has the property that it's gradient can be expressed in terms of the policy. This also means, that we can express the advantage function by taking the derivative w.r.t. the policy and multiplying it by w :

$$\nabla_w A_w(s, a) = \nabla_\theta \log \pi_\theta(s|a) \quad (12)$$

$$A_w(s, a) = \nabla_\theta \log \pi_\theta(s|a) w. \quad (13)$$

Inserting this and bringing the cost function $J(\theta)$ to the left hand side:

$$\sum_{i=0}^N \gamma^i \nabla_\theta \log \pi_\theta(a_i|s_i)^T \cdot w + J(\theta) = \sum_{i=0}^N \gamma^i r(s_i, a_i) \quad (14)$$

This is exactly the equation, which we solve in algorithm 1 by taking the left pseudo inverse. Besides the parameter vector w , we receive the cost function $J(\theta)$ as a side product.

The update of the policy parameters in direction of the critic parameters can

also be explained by the compatible function approximation (equation 13). With this the natural policy gradient [REF: EQUATION] simplifies:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s|a) A_w(s, a)] \quad (15)$$

$$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s|a) \nabla_{\theta} \log \pi_{\theta}(s|a)^T w] \quad (16)$$

$$= G_{\theta} w \quad (17)$$

$$\tilde{\nabla}_{\theta} J(\theta) = w \quad (18)$$

We show several modifications to the algorithm in the next section and discuss its advantages and drawbacks.

5 NAC Modifications

Least Squares: Besides the episodic NAC, Peters specified another approach, NAC using LSTD-Q(λ) [17, 7], in his original Paper [26]. The main difference is the estimation of the critic’s parameters. LSTD-Q(λ) uses least squares temporal difference learning [6] to estimate the parameters of the critic after every step taken in the environment. The algorithm uses eligibility traces [29] and two linear functions: $A_w(s, a) = \nabla_{\theta} \log \pi_{\theta}(a|s)^T w$ and $V_v(s) = \phi(s)v$. The latter is an approximation of the value function and needed to update the critic by solving a linear set of equations induced by the least squares error between our observation and a value function approximation. As in the episodic case, when the angle between the critics parameters is smaller than a fixed value, the actor is updated a proportion in the direction of the critic.

Recursive Least Squares:: The “RLS-based natural actor-critic algorithm” equips the LSTD-Q NAC algorithm with a recursive update rule for the parameters of the critic [20]. The old parameter values are reused during the current update which increases efficiency [36].

Fitted NAC & Importance Sampling: Fitted natural actor critic (FNAC) is a fitted version of the natural actor critic algorithm [18]. It employs a memory D , which is filled by sampling data from the environment. Once filled, the least squares NAC algorithm is executed as usual. Normally, after an policy update, we would need to sample D again, this time with the improved policy. However, Melo et al. implemented importance sampling (IS) [29] to avoid the re-sampling. In addition to the current policy parameters θ , IS also saves the policy parameters θ^- , which were used to sample the memory D . Every time we evaluate the critic, we multiply our estimation by the importance weights, $\frac{\pi_{\theta}(a|s)}{\pi_{\theta^-}(a|s)}$, to estimate the proportion we need to change the current critic parameters. The memory D is independent of the current learning policy. This approach is extremely data efficient and brings fundamental advantages in situations, where collecting data is costly or time consuming. Additionally, FNAC makes use of regression methods to update the critic’s parameters,

which allow the use of a general function approximator for the value function instead of compulsory linear one. This positively impacts the accuracy of the critic’s estimation.

Incremental NAC (INAC): The incremental NAC algorithm combines linear function approximation and bootstrapping [5]. It reuses existing approach, namely temporal difference learning [29] and two-timescale stochastic approximation [4]. Bhatnagar et al. provide three new natural gradient algorithms and proved that they locally converge to a local maximum. The main feature of the algorithms is the totally incremental estimation of the policy, the policy is changed every time step, and the incrementally update of the gradient. In comparison to Peters et al., the policy gradient is not reset every update, but saved and reused to support calculating the gradient of the next iteration. These improvements facilitate the application to large-scale reinforcement learning problems, decreases computation time and makes the algorithm more efficient than conventional actor critic methods. Further, one of the algorithms can be executed without explicitly computing the inverse Fischer information matrix. This authors report even faster convergence.

Implicit Incremental NAC (I2NAC): INAC algorithms suffer from a difficult to tune step size and an unstable and sometimes divergent estimation of the natural gradient. Iwaki et al. analyzed the reasons for these drawbacks and created an improved algorithm, the implicit incremental NAC algorithm [13]. This algorithm uses the ideas of implicit stochastic gradient descent [33] and the implicit temporal differences [31] to overcome these difficulties. The change between INAC and I2NAC is a weight vector, which is multiplied with the update of the critic’s parameter vector. It makes use of the eligibility traces and a new hyper parameter β . This stabilizes learning and empirical results show less divergence.

Regularization on NAC: Even if we find the inverse of G , it can be ill defined. An example for this are extremely small eigenvalues which appear due to noise in our data. These eigenvalues will become extremely large if we take the inverse of G and thus the parameters belonging to the eigenvalues will get a lot of credibility which they should not have and which will falsify our inverse.

That is why there have been some approaches to introduce a regularization term [28]. Regularizing the matrix inverse can for example be done by a technique called stochastic robust approximation [8], where G^{-1} is replaced by $G_{\text{reg}}^{-1} = (G^T G + \epsilon I)^{-1} G^T$ and ϵ denotes a small constant (e.g 0.01).

Witsch et al. use an approach similar to least squares regularization [35]: $\tilde{\nabla}_{\theta} J(\theta) = (F + \lambda I)^{-1} \nabla_{\theta} J(\theta)$. If λ is huge, the Fisher matrix only has a small influence on the change in direction. Therefore, we want to scale λ regarding

F : $\lambda = \frac{\alpha}{\det(\tilde{F})+1}$, where α is small constant, e.g. 0.01.

Another idea is the application of ridge regression [11], which has a build in regularizer. We can calculate $\tilde{\nabla}_\theta J(\theta)$ by solving the linear equation $G(\theta)\tilde{\nabla}_\theta J(\theta) = \nabla_\theta J(\theta)$ in the direction of $\tilde{\nabla}_\theta J(\theta)$.

POMDPs: There have been first approaches to apply the NAC to POMDPs. A promising approach is the Natural Actor and Belief Critic [14], which modifies NAC to learn parameters using belief states in statistical dialogue systems.

6 Discussion

In this paper we described the natural gradient, the natural actor critic algorithm and modifications which have been applied to the NAC in the last years. NAC is a state of the art algorithm which can be applied model-free and with continuous action spaces. It has been reported, that NAC converges faster than vanilla gradient methods and that it can jump out of plateaus. These are key advantages why we expect to see more use of NAC's in the future.

Disadvantages are clearly the need to invert a matrix and the extinction of variance. For the first, people may need to come up with a solution to faster invert matrices or even how to avoid the estimation. The latter has already been tackled by some algorithms, namely TRPO and PPO. A very alarming study claims that NAC methods exhibit, in contrast to the expectations, a bias [32]. Further research have to shed light on this issue.

Further, some of the modifications still need application to real-world problems to assess their ultimate utility. Most of the times the modifications are only applied in a special way and further research is needed to incorporate eligibility traces, least-squares methods, online implementation extend it to the deterministic policy gradient method.

References

1. Shun-ichi Amari. Differential geometry of a parametric family of invertible linear systemsriemannian metric, dual affine connections, and divergence. *Mathematical systems theory*, 20(1):53–82, 1987.
2. Shun-Ichi Amari. Natural gradient works efficiently in learning. *Neural computation*, 10(2):251–276, 1998.
3. Shun-Ichi Amari and Scott C Douglas. Why natural gradient? In *Proceedings of the 1998 IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP’98 (Cat. No. 98CH36181)*, volume 2, pages 1213–1216. IEEE, 1998.
4. Shalabh Bhatnagar and Vivek S Borkar. A two timescale stochastic approximation scheme for simulation-based parametric optimization. *Probability in the Engineering and Informational Sciences*, 12(4):519–531, 1998.
5. Shalabh Bhatnagar, Mohammad Ghavamzadeh, Mark Lee, and Richard S Sutton. Incremental natural actor-critic algorithms. In *Advances in neural information processing systems*, pages 105–112, 2008.
6. Justin A Boyan. Least-squares temporal difference learning. In *ICML*, pages 49–56, 1999.
7. Justin A Boyan. Technical update: Least-squares temporal difference learning. *Machine learning*, 49(2-3):233–246, 2002.
8. Stephen Boyd and Lieven Vandenberghe. *Convex optimization*. Cambridge university press, 2004.
9. Guillaume Desjardins, Razvan Pascanu, Aaron Courville, and Yoshua Bengio. Metric-free natural gradient for joint-training of boltzmann machines. *arXiv preprint arXiv:1301.3545*, 2013.
10. Simon S Haykin, Simon S Haykin, Simon S Haykin, Kanada Elektrotechnik, and Simon S Haykin. *Neural networks and learning machines*, volume 3. Pearson Upper Saddle River, 2009.
11. Arthur E Hoerl and Robert W Kennard. Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics*, 12(1):55–67, 1970.
12. Ronald A Howard. Dynamic programming and markov processes. 1960.
13. Ryo Iwaki and Minoru Asada. Implicit incremental natural actor critic algorithm. *Neural Networks*, 109:103–112, 2019.
14. Filip Jurčiček, Blaise Thomson, and Steve Young. Natural actor and belief critic: Reinforcement algorithm for learning parameters of dialogue systems modelled as pomdps. *ACM Transactions on Speech and Language Processing (TSLP)*, 7(3):6, 2011.
15. Sham M Kakade. A natural policy gradient. In *Advances in neural information processing systems*, pages 1531–1538, 2002.
16. Byungchan Kim, Jooyoung Park, Shinsuk Park, and Sungchul Kang. Impedance learning for robotic contact tasks using natural actor-critic algorithm. *IEEE Transactions on Systems, Man, and Cybernetics, Part B (Cybernetics)*, 40(2):433–443, 2010.
17. Michail G Lagoudakis and Ronald Parr. Least-squares policy iteration. *Journal of machine learning research*, 4(Dec):1107–1149, 2003.

18. Francisco S Melo and Manuel Lopes. Fitted natural actor-critic: A new algorithm for continuous state-action mdps. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, pages 66–81. Springer, 2008.
19. Arkadi Nemirovski. Efficient methods in convex programming. 2005.
20. Jooyoung Park, Jongho Kim, and Daesung Kang. An rls-based natural actor-critic algorithm for locomotion of a two-linked robot arm. In *International Conference on Computational and Information Science*, pages 65–72. Springer, 2005.
21. Razvan Pascanu and Yoshua Bengio. Revisiting natural gradient for deep networks. *arXiv preprint arXiv:1301.3584*, 2013.
22. Jan Peters and Stefan Schaal. Applying the episodic natural actor-critic architecture to motor primitive learning. In *ESANN*, pages 295–300, 2007.
23. Jan Peters and Stefan Schaal. Natural actor-critic. *Neurocomputing*, 71(7-9):1180–1190, 2008.
24. Jan Peters and Stefan Schaal. Reinforcement learning of motor skills with policy gradients. *Neural networks*, 21(4):682–697, 2008.
25. Jan Peters, Sethu Vijayakumar, and Stefan Schaal. Reinforcement learning for humanoid robotics. In *Proceedings of the third IEEE-RAS international conference on humanoid robots*, pages 1–20, 2003.
26. Jan Peters, Sethu Vijayakumar, and Stefan Schaal. Natural actor-critic. In *European Conference on Machine Learning*, pages 280–291. Springer, 2005.
27. Silvia Richter, Douglas Aberdeen, and Jin Yu. Natural actor-critic for road traffic optimisation. In *Advances in neural information processing systems*, pages 1169–1176, 2007.
28. Jascha Sohl-Dickstein. The natural gradient by analogy to signal whitening, and recipes and tricks for its use. *arXiv preprint arXiv:1205.1828*, 2012.
29. Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*. MIT press, 2018.
30. Richard S Sutton, David A McAllester, Satinder P Singh, and Yishay Mansour. Policy gradient methods for reinforcement learning with function approximation. In *Advances in neural information processing systems*, pages 1057–1063, 2000.
31. Aviv Tamar, Panos Toulis, Shie Mannor, and Edoardo M Airolidi. Implicit temporal differences. *arXiv preprint arXiv:1412.6734*, 2014.
32. Philip Thomas. Bias in natural actor-critic algorithms. In *International Conference on Machine Learning*, pages 441–448, 2014.
33. Panagiotis Toulis, Edoardo Airolidi, and Jason Rennie. Statistical analysis of stochastic gradient methods for generalized linear models. In *International Conference on Machine Learning*, pages 667–675, 2014.
34. Ronald J Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning*, 8(3-4):229–256, 1992.

35. Andreas Witsch, Roland Reichle, Kurt Geihs, Sascha Lange, and Martin Riedmiller. Enhancing the episodic natural actor-critic algorithm by a regularisation term to stabilize learning of control structures. In *2011 IEEE Symposium on Adaptive Dynamic Programming and Reinforcement Learning (ADPRL)*, pages 156–163. IEEE, 2011.
36. Xin Xu, Han-gen He, and Dewen Hu. Efficient reinforcement learning using recursive least-squares methods. *Journal of Artificial Intelligence Research*, 16:259–292, 2002.