

Natural Actor Critic

Components and Extensions

Maximilian Gehrke

Tabea Wilke

Yannik Frisch

Group 19 Oleg Arenz



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- Optimization problem:

$$\begin{aligned} \max_{\delta\theta} J(\theta + \delta\theta) &\approx J(\theta) + \delta\theta^T \nabla_{\theta} J(\theta) \\ \text{s.t. } \epsilon = D_{KL}(\pi_{\theta} || \pi_{\theta+\delta\theta}) &\approx \frac{1}{2} \delta\theta^T F_{\theta} \delta\theta \end{aligned}$$

- Solution:

$$\tilde{\nabla}_{\theta} J(\theta) = F_{\theta}^{-1} \nabla_{\theta} J(\theta)$$

- Fisher Information Matrix:

$$F_{\theta} = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) \nabla_{\theta} \log \pi_{\theta}(a|s)^T]$$

⇒ *Parametrization invariant, data efficient & fast convergence*

The Natural Actor Critic algorithm



Algorithm 1 Episodic Natural Actor Critic (eNAC)

Require: Parameterized policy $\pi_\theta(a|s)$ and its derivative $\nabla_\theta \log \pi_\theta(a|s)$
with initial parameters $\theta = \theta_0$.

- 1: **for** $u = 1, 2, 3, \dots$ **do**
 - 2: **for** $e = 1, 2, 3, \dots$ **do**
 - 3: **Execute roll-out:** Draw initial state $s_0 \sim p(s_0)$
 - 4: **for** $t = 1, 2, 3, \dots, N$ **do**
 - 5: Draw action $a_t \sim \pi_{\theta_t}(a_t|s_t)$, observe next state $s_{t+1} \sim p(s_{t+1}|s_t, a_t)$
and reward $r_{t+1} = r(s_t, a_t)$.
 - 6: **end for**
 - 7: **end for**
 - 8: **Critic Evaluation (repeat for each sampled trajectory):** Determine compatible
function approximation of advantage function $A(s, a) \approx A_{w_t}(s, a)$.
 - 9: Determine basis functions: $\Phi_e = \left[\sum_{t=0}^T \gamma^t \nabla_\theta \log \pi_\theta(a_t|s_t)^T, 1 \right]^T$,
reward statistics: $R_e = \sum_{t=0}^T \gamma^t r_t$ and solve $\begin{bmatrix} w_e \\ J \end{bmatrix} = (\Phi_e^T \Phi_e)^{-1} \Phi_e^T R_e$.
Update critic parameters: $w_{t+1} = w_t + \beta w_e$.
 - 10: **Actor Update:** When the natural gradient is converged, $\angle(w_{t+1}, w_t) \leq \epsilon$, update
the policy parameters: $\theta_{t+1} = \theta_t + \alpha w_{t+1}$.
 - 11: **end for**
-

Extensions



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- ▶ Recursive Least Squares
- ▶ Fitted NAC + Importance Sampling (FNAC)
- ▶ Incremental NAC (INAC)
- ▶ Implicit Incremental NAC (I2NAC)
- ▶ Regularization



- ▶ NAC has several advantages over vanilla PGM
- ▶ Different approaches for:
 - ▶ Updating critic (& adjusting learning rate)
 - ▶ Fisher inverse
 - ▶ Actor update frequency
- ▶ Open Questions:
 - ▶ Inverse of Fisher Information Matrix (expensive!)
 - ▶ NPG estimation might be biased (Thomas, 2014)
 - ▶ POMDP's (Jurčíček u. a., 2011)

For publication references please see our paper “*Natural Actor Critic: Components and Extensions*”.

- [Jurčiček u. a. 2011] JURČÍČEK, Filip ; THOMSON, Blaise ; YOUNG, Steve:
Natural actor and belief critic: Reinforcement algorithm for learning parameters
of dialogue systems modelled as POMDPs. In: *ACM Transactions on Speech
and Language Processing (TSLP)* 7 (2011), Nr. 3, S. 6
- [Thomas 2014] THOMAS, Philip: Bias in natural actor-critic algorithms. In:
International Conference on Machine Learning, 2014, S. 441–448